



Assignment Title:

Dynamic Robotics Simulation and Autonomous System Design

Programme title:

MSc Artificial Intelligence

Prepared by:

Amaan Mohammed Khalander

Module:

Robotics Software and Programming

Year:

Term 2 – 2026

Submission Date:

9 / October / 2026

Issued by:

Dr. Giulio Napolitano, Dr. Syed Arslan Abbas Rizvi

Learner declaration

Item	Response
Word count	3,137 (Main body including headings and code, excluding captions and tables)
Use of proof-reader/proof-reading service	NO
If yes, specify which service	N/A
I confirm that I submit this work as my own work and have cited all sources used	YES
I confirm that I have followed guidance on the acceptable use of AI tools for this assignment	YES
Where use of appropriately cited AI tools is permitted, I confirm that I have cited in accordance with the UCA Harvard Referencing Standard	YES

Student signature	Date
	9 / October / 2026

Source Code and Relevant Project Links:

Source Code:

- Task 1:
<https://github.com/Amaan9136/bsbi-assignments/tree/main/Sem2/Robotics%20Software%20and%20Programming/task1-code>
- Task 2:
<https://github.com/Amaan9136/bsbi-assignments/tree/main/Sem2/Robotics%20Software%20and%20Programming/task2-code>

Table of Contents

Abstract	6
Introduction.....	7
Task 1: Interactive PyBullet Environment with Object Semantics.....	8
3.1 Project Definition.....	8
3.2 Literature Review.....	8
3.3 Design and Implementation	8
3.4 Code Implementation.....	10
3.5 Execution and Testing.....	11
3.6 Evaluation and Reflection.....	12
Task 2: Goal-Oriented Autonomous Robot with Monitoring (ROS2).....	14
4.1 Project Definition.....	14
4.2 Design and Implementation	14
4.3 Execution and Testing.....	17
4.4 Evaluation and Reflection.....	22
Combined Evaluation and Reflection	23
Concluding Remarks.....	24
Acknowledgement	25
References.....	26

Table of Figures

Figure 1: Task 1 pipeline, from the floor-plan sketch and URDF generation through the semantic layer and camera rendering to the query helpers and outputs.	9
Figure 2: Six PyBullet camera viewpoints, robot repositioned between renders while all other objects remain fixed.	12
Figure 3: ROS2 node and topic architecture of the semantic_nav_monitor package, showing operator input, the four nodes and the Gazebo simulation.	15
Figure 4: Six-state finite state machine in mission_controller, with the REVERSE, ORIENT and COMMIT sub-phases shown beneath their parent states.	16
Figure 5: TheConstruct.ai Rosject XctZpczWOA: file explorer, mission_controller.py in the code editor, the live Simulation panel, and monitor_node's summary stream in the shell.	18
Figure 6: Gazebo client on TheConstruct.ai: robot near the charging dock among the eight pallet obstacles and checkpoint markers, with elapsed real time and iteration count visible. .	19
Figure 7: keyboard_hri_node in a Rosject shell: start and pause keypresses published and confirmed, alongside the live Simulation panel.	19
Figure 8: Live /mission_state echo showing the robot cycle through AVOID_OBSTACLE, REPLAN, IDLE, PLANNING, NAVIGATE and MISSION_COMPLETE.	20
Figure 9: monitor_node's mission summary, captured while the mission was in progress.	20
Figure 10: Local Gazebo Harmonic verification: LIDAR rays visualised around the robot near a pallet obstacle, with the forward camera feed shown alongside.	22

List of Tables

Table 1: Semantic object registry as printed by print_object_roles() (14 objects, excluding the robot).	10
Table 2: ROS2 node architecture.	15
Table 3: Finite state machine (six states).	16
Table 4: Mission summary logged by monitor_node on TheConstruct.ai, Rosject XctZpczWOA, while the mission was in progress.	21
Table 5: Task 2 test cases and observed outcomes.	21

Statement of compliance with academic ethics and the avoidance of plagiarism

I honestly declare that this dissertation is entirely my own work, and that none of its parts has been copied from printed or electronic sources, translated from foreign sources, or reproduced from essays by other researchers or students. Wherever I have based my work on ideas or other people's texts, I clearly declare it through proper references in line with academic ethics.

(In the case that is proved that part of the essay does not constitute an original work, but a copy of an already published essay or from another source, the student will be expelled permanently from the postgraduate program).

Name and Surname:

AMAAN MOHAMMED KHALANDER

A handwritten signature in black ink, appearing to read 'Amaan MK', with a stylized flourish at the end.

Date: 9 / October / 2026

Abstract

This report covers two exercises for the Robotics Software and Programming module. Task 1 is a PyBullet simulation of an inspection and light-warehousing room with five URDF object categories, fourteen objects, a two-wheel robot and six camera viewpoints. A dictionary called `semantic_db` stores each object's role for lookups. Task 2 is a TurtleBot3 patrol robot in ROS2 with a six-state finite state machine. The world is a custom Gazebo warehouse with eight pallet obstacles and four checkpoints. One checkpoint is the charging dock. The robot plans its route first. It uses LIDAR and odometry to avoid obstacles and follows keyboard commands. A monitoring node logs state transitions and mission metrics. The mission reached `MISSION_COMPLETE` on TheConstruct.ai in Rosject XctZpczWOA, and a local ROS2 Jazzy and Gazebo Harmonic workspace was used for extra checks.

Introduction

The module asks us to design, build and test robot software in simulation, with no physical hardware. The brief, Dynamic Robotics Simulation and Autonomous System Design, sets out two stages of increasing complexity. Task 1 is a PyBullet room where every object has a role. Task 2 is a ROS2 robot that senses its surroundings, navigates, takes keyboard input and logs its own performance. The tasks use different tools, but the guidelines ask for one report. Both are in this report, in the BSBI template.

In Task 1, each object in the room has a role (obstacle, target, landmark, shelf or charger). The role is stored in code and can be looked up. Task 2 moves from a fixed scene to a moving robot. It reads LIDAR and odometry, switches between states and logs what it did and why.

Task 1: Interactive PyBullet Environment with Object Semantics

3.1 Project Definition

The objective of Task 1 was a PyBullet scene where code could look up each object's role as well as its position. I sketched the room on paper before writing any code. six by four metres with one entrance, a central aisle, back-wall shelving, a charging dock near the entrance and a landmark at each corner. I then turned the sketch into URDF categories and coordinates.

The finished scene has fourteen objects in five URDF categories, which is above the brief's minimum of eight objects and four models. A two-wheel robot is moved between camera renders to show movement. All objects are URDF boxes and cylinders with no imported meshes, so the notebook is self-contained.

3.2 Literature Review

Three areas of work informed Task 1, covering simulators, semantic representation and labelled scenes.

- **Simulators.** PyBullet gives Python access to the Bullet physics engine. It loads URDF models and has a CPU renderer (Coumans and Bai, 2021; PyBullet, 2023), so the notebook runs headless in Colab. Collins et al. (2021) reviewed the physics simulators used in robotics. Task 1 only needs a scene and some renders, so PyBullet was enough.
- **Semantic representation.** Kostavelis and Gasteratos (2015) survey semantic mapping for mobile robots, noting its role in human-robot interaction, knowledge representation and planning. The registry in Task 1 is a very simple version of that idea. I assign the labels by hand, so code can ask what an object is as well as where it is.
- **Labelled scenes.** McCormac et al. (2017) generated SceneNet RGB-D, a large synthetic indoor dataset with ground-truth labels. Task 1 is a small hand-made scene. It is less realistic, but I control every label.

3.3 Design and Implementation

The notebook (Task1_PyBullet_Object_Semantics.ipynb) is written in Python 3 and has nine executable cells. The setup cell installs pybullet if it is missing, then imports pybullet, pybullet_data, numpy and matplotlib. The generation cell writes six URDF files, one each for the obstacle (crate), target zone, landmark, shelf and charging dock, plus the two-wheel robot with a box chassis and two cylindrical wheels. The environment cell connects to PyBullet in DIRECT mode, loads a ground plane from pybullet_data and places the fourteen objects and

the robot at the coordinates from my floor sketch. The scene has four crates beside the aisle, three target zones (pickup, an off-aisle inspection point and drop-off), four corner landmarks, two back-wall shelves and one charging dock. Figure 1 shows the steps in order.




Figure 1: Task 1 pipeline, from the floor-plan sketch and URDF generation through the semantic layer and camera rendering to the query helpers and outputs.

The semantic layer, `semantic_db`, is a plain Python dictionary. Each PyBullet body id maps to a role string and an (x, y, z) position. Table 1 shows it as printed. Three helper functions use it. `get_objects_by_role` filters by role string. `get_nearest_target` measures the Euclidean distance from the robot's current base position to every target and returns the closest. `print_object_roles` prints the whole registry.

The rendering cell sets up six camera viewpoints with `computeViewMatrix` and `computeProjectionMatrixFOV` (PyBullet, 2023). The brief requires three, and I rendered six as an extension. Only the robot moves between renders. Every other object stays fixed. `getCameraImage` returns a 480 by 360 RGB array for each view.

The tools were Python 3, PyBullet, NumPy and Matplotlib. The notebook has no perception model. Roles are set when the scene is built, so it cannot work out a role from an image (see Section 3.6).

Table 1: Semantic object registry as printed by `print_object_roles()` (14 objects, excluding the robot).

body_id	role	position (x, y, z) m
1	obstacle	(1.50, 0.80, 0.25)
2	obstacle	(1.50, -0.80, 0.25)
3	obstacle	(3.50, 0.80, 0.25)
4	obstacle	(3.50, -0.80, 0.25)
5	target	(0.50, 0.00, 0.01)
6	target	(3.00, 1.60, 0.01)
7	target	(5.50, 0.00, 0.01)
8	landmark	(0.30, 1.80, 0.70)
9	landmark	(0.30, -1.80, 0.70)
10	landmark	(5.70, 1.80, 0.70)
11	landmark	(5.70, -1.80, 0.70)
12	shelf	(5.50, 1.00, 0.80)
13	shelf	(5.50, -1.00, 0.80)
14	charger	(0.30, 0.90, 0.07)

3.4 Code Implementation

The system is a Python Jupyter notebook. It runs locally or in Colab. The excerpts below show the main parts.

URDF generation (cell 3): each category is a short URDF string written to disk; the obstacle is shown.

```
OBSTACLE_URDF = """<?xml version="1.0"?>
<robot name="obstacle">
  <link name="base_link">
    <visual><geometry>
      <box size="0.5 0.5 0.5"/></geometry></visual>
    <collision><geometry>
      <box size="0.5 0.5 0.5"/></geometry></collision>
    <inertial><mass value="0"/></inertial>
  </link>
</robot>"""
URDF_PATHS["obstacle"] = write_urdf("obstacle.urdf", OBSTACLE_URDF)
```

Semantic layer and query (cell 5): `semantic_db` maps body ids to roles; `get_nearest_target` reads the robot's live pose and returns the closest target.

```

semantic_db = {}
for body_id, role, pos in object_ids:
    semantic_db[body_id] = {"role": role, "position": pos}

def get_nearest_target(robot_body_id):
    robot_pos, _ = p.getBasePositionAndOrientation(robot_body_id)
    targets = get_objects_by_role("target")
    best, best_dist = None, float("inf")
    for body_id, pos in targets:
        dist = math.dist(robot_pos, pos)
        if dist < best_dist:
            best_dist, best = dist, (body_id, pos)
    return {"body_id": best[0], "position": best[1],
            "distance": best_dist}

```

Camera rendering (cell 6): fixed eye and target pairs are rendered on the CPU in DIRECT mode; only the robot's pose changes between calls.

```

def render_view(camera_eye, camera_target, up_vector=(0, 0, 1)):
    view = p.computeViewMatrix(camera_eye, camera_target, up_vector)
    proj = p.computeProjectionMatrixFOV(
        fov=60, aspect=IMG_WIDTH / IMG_HEIGHT,
        nearVal=NEAR_PLANE, farVal=FAR_PLANE)
    _, _, rgb, _, _ = p.getCameraImage(
        IMG_WIDTH, IMG_HEIGHT, view, proj,
        renderer=p.ER_TINY_RENDERER)
    return np.reshape(rgb, (IMG_HEIGHT, IMG_WIDTH, 4))[:, :, :3]

```

3.5 Execution and Testing

I ran the notebook top to bottom in a local Python 3.12 environment with PyBullet 3.2.7, and then once more in Google Colab. The completed notebook is hosted publicly on GitHub at github.com/Amaan9136/bsbi-assignments, under Sem2/Robotics Software and Programming/task1-code/Task1_PyBullet_Object_Semantics.ipynb, and opens in Google Colab via https://colab.research.google.com/github/Amaan9136/bsbi-assignments/blob/main/Sem2/Robotics%20Software%20and%20Programming/task1-code/Task1_PyBullet_Object_Semantics.ipynb. All nine cells ran without errors.

Figure 2 shows the six rendered viewpoints and the robot pose in each. The cameras are fixed. Viewpoint 1 looks down on the whole room, and viewpoints 2 to 6 are oblique views from different sides. The robot stands at (0.0, 0.0), (3.0, 0.0), (5.5, 0.0), (1.5, 0.15), (3.0, 1.6) and (5.0, -1.5) in viewpoints 1 to 6. The overhead view shows every object and the robot. The other views show the same objects from the side.




Figure 2: Six PyBullet camera viewpoints, robot repositioned between renders while all other objects remain fixed.

In cell 8 I moved the robot to (5.5, 0.0, 0.1), next to the drop-off target, and ran the three helper functions. `print_object_roles` printed the fourteen-row registry in Table 1. `get_objects_by_role("obstacle")` returned the four crates (ids 1 to 4). `get_objects_by_role("shelf")` returned the two shelves (ids 12 and 13). `get_nearest_target` returned body id 7, the drop-off target at (5.50, 0.00, 0.01), 0.092 m away.

3.6 Evaluation and Reflection

The hardest part was keeping the camera framing steady while the robot moved. My first version worked out each camera position from the robot's pose. The framing kept shifting, and in one configuration the robot left the frame completely. I fixed the cameras at positions based on the room bounds and moved only the robot between renders. I began with the three views the brief asks for, then went to six so I could check the objects from more angles. I also tried matching URDF filenames to role strings when loading objects. I dropped it because it was easy to get wrong. Setting each role in `SCENE_LAYOUT` keeps all roles in one place, so they are easy to check.

The main limitation is that I set the labels by hand. The robot cannot recognise a role from how an object looks. The brief asks for an interpretable environment and not a perception system,

so this was acceptable for Task 1. I could train a small classifier on SceneNet RGB-D (McCormac et al., 2017) to predict roles from the Section 3.5 images. In a real warehouse, a wrong role could make the robot hit an obstacle or ignore a target. So I kept the link between geometry and role clear.

Task 2: Goal-Oriented Autonomous Robot with Monitoring (ROS2)

4.1 Project Definition

Task 2 uses a warehouse inspection patrol robot. A TurtleBot3 starts at the charging dock and drives a loop of four checkpoints around a warehouse bay. It avoids pallets on the way, takes simple operator commands and reports its own performance. I chose it because one mission covers navigation, obstacle avoidance, keyboard control and monitoring.

The package, `semantic_nav_monitor`, is built for TheConstruct.ai, which runs ROS2 Humble (Open-Source Robotics Foundation, 2024) with Gazebo Classic and the standard TurtleBot3 stack (ROBOTIS, 2023). I built and tested it on TheConstruct.ai (TheConstruct, 2023) in Rosject `XctZpczWOA` (URL identifier 1051220, `app.theconstruct.ai/desktop/rosject/1051220`), workspace `ros2_ws`. The free tier gives no permanent shareable URL, so I made the project public under the title `XctZpczWOA`, the unique name the brief asks for. I also used a local Ubuntu 24.04 workspace with ROS2 Jazzy and Gazebo Harmonic (Open Robotics, 2024) to view the LIDAR rays, camera feed and planned-path toggle outside the browser. The GitHub folder holds this Jazzy and Harmonic version. The Rosject copy uses a Gazebo Classic launch file and world.

The custom Gazebo world, `warehouse_inspection`, is a walled bay of 7 m by 7 m. It comes as a Classic `.world` for TheConstruct.ai and a Harmonic `.sdf` for local checks. It contains eight pallet stacks (two per patrol leg), three small pylons on the first leg, three coloured checkpoint markers and a green charging dock marker at the origin. The robot is a TurtleBot3 Burger with a forward camera (`burger_cam`). Its 360-degree LIDAR publishes `/scan` and its wheel odometry publishes `/odom`.

4.2 Design and Implementation

Two bodies of published work shaped the behaviour design. Colledanchise and Ögren (2018) describe finite state machines as becoming brittle as a system grows, which is why behaviour trees are now common in larger robot software. This patrol mission has one goal and a handful of states, so I kept a finite state machine because it is easier to read and to log. For avoidance, the Bug algorithms and the vector field histogram steer from local sensor data (Siegwart, Nourbakhsh and Scaramuzza, 2011; Borenstein and Koren, 1991). My escape-heading probe is a much simpler relative of that idea.

The package contains four ROS2 nodes (Quigley, Gerkey and Smart, 2015), summarised in Table 2. `mission_controller` is the only node that commands the robot. `monitor_node` only observes. `keyboard_hri_node` reads single keypresses without needing Enter. `path_visualizer_node` draws the planned route in the Gazebo Sim client and does not affect the robot's behaviour. It needs the `gz` command line tool, so it works only in the Harmonic build. Figure 3 shows the node and topic architecture.




Figure 3: ROS2 node and topic architecture of the `semantic_nav_monitor` package, showing operator input, the four nodes and the Gazebo simulation.

Table 2: ROS2 node architecture.

Node	Subscribes	Publishes	Role
<code>mission_controller</code>	<code>/scan</code> , <code>/odom</code> , <code>/hri_command</code>	<code>/cmd_vel</code> , <code>/mission_state</code> , <code>/planned_path</code>	Six-state FSM: navigation, obstacle avoidance, replanning
<code>monitor_node</code>	<code>/mission_state</code> , <code>/odom</code>	(console log only)	Logs transitions and computes time-per- state, distance and obstacle count

Node	Subscribes	Publishes	Role
keyboard_hri_node	terminal keypresses	/hri_command, /show_planned_path	Single-key HRI: s start, p pause, x stop, v toggle path, q quit
path_visualizer_node	/planned_path, /show_planned_path	spawns/removes Gazebo line entities	Draws or hides the planned route in the Gazebo client (local Harmonic build only)

mission_controller implements six states. The brief's minimum is five, and I added PLANNING ahead of NAVIGATE. Table 3 lists each state and Figure 4 shows the transitions.




Figure 4: Six-state finite state machine in mission_controller, with the REVERSE, ORIENT and COMMIT sub-phases shown beneath their parent states.

Table 3: Finite state machine (six states).

State	Entered when	Behaviour
IDLE	At launch, or after the mission aborts	Stationary, waiting for a start command
PLANNING	start received	Builds the four-checkpoint route as a nav_msgs/Path, publishes it once, then hands off immediately
NAVIGATE	After planning, or once an obstacle is cleared	Proportional controller drives towards the next checkpoint
AVOID_OBSTACLE	Forward clearance below 0.32 m for three consecutive control	Reverses if closer than 0.20 m (REVERSE), then rotates to a

State	Entered when	Behaviour
	cycles, or an odometry-based stall	LIDAR-probed clear heading (ORIENT)
REPLAN	A clear heading has been reached	Drives forward roughly 0.45 m along that heading (COMMIT) before resuming NAVIGATE
MISSION_COMPLETE	All four checkpoints visited, or stop received	Robot halts and final state published

I made two design decisions. First, obstacle avoidance has three phases (REVERSE, ORIENT and COMMIT). The robot probes escape headings at eleven angles per side using LIDAR and scores each by how well it lines up with the goal. Among similar scores it takes the heading with more open space. A heading is accepted only with at least 0.21 m of clear runway, which is the robot's half-width (0.11 m) plus a 0.05 m lateral margin and a further 0.05 m. An odometry stall check runs alongside. If the robot moves less than 0.05 m in one second of commanded forward motion, it counts as blocked whatever the LIDAR says. After ten failed attempts the mission aborts to IDLE. Second, in the Jazzy and Harmonic build the velocity commands are TwistStamped, because the `ros_gz_bridge` started by the TurtleBot3 launch files expects that type. An extra bridge in an earlier launch file subscribed to `/cmd_vel` as plain Twist, and the mismatch stopped the robot moving until I removed that bridge.

`monitor_node` subscribes to `/mission_state` and `/odom`. On every state change it logs the transition with a timestamp and adds the elapsed time to the previous state. Each entry into `AVOID_OBSTACLE` logs "Obstacle detected -> switching to AvoidObstacle" and increments an obstacle-encounter counter, which stands in for the brief's collisions metric since the robot avoids contact. `mission_controller` also logs messages such as "Reached Checkpoint 1 - East Aisle". Odometry feeds a running distance total. A five-second timer prints a summary during the mission, and a final one is printed on `MISSION_COMPLETE`. Times use the system clock, so they are wall-clock seconds. `path_visualizer_node` spawns thin static box entities along the route through the Gazebo EntityFactory service, and the 'v' key or a native GUI panel toggles the line without changing the plan.

4.3 Execution and Testing

In the Rosject I ran `cd ~/ros2_ws,` then `colcon build --packages-select semantic_nav_monitor,` then `source install/setup.bash.` I set

TURTLEBOT3_MODEL=burger_cam, then ran `ros2 launch semantic_nav_monitor semantic_nav_monitor.launch.py` in Shell 1 and `ros2 run semantic_nav_monitor keyboard_hri_node` in Shell 2. A notebook inside the Rosject lists these steps and every shell command.

Figure 5 shows the Rosject workspace. The file explorer shows the name XctZpczWOA, `mission_controller.py` is open in the code editor and the Simulation panel shows the warehouse world. A shell is printing `monitor_node`'s summary.




Figure 5: TheConstruct.ai Rosject XctZpczWOA: file explorer, `mission_controller.py` in the code editor, the live Simulation panel, and `monitor_node`'s summary stream in the shell.

Figure 6 shows the Gazebo client during the run. The status bar shows 00:02:11 of real time, 94,858 iterations and a real-time factor of 0.70, so the simulation was running.




Figure 6: Gazebo client on TheConstruct.ai: robot near the charging dock among the eight pallet obstacles and checkpoint markers, with elapsed real time and iteration count visible.

Figure 7 shows keyboard_hri_node in a Rosject shell next to the Simulation panel. The s and p keypresses are echoed as Published HRI command: 'start' and 'pause', and Planned-path line: ON confirms that the path toggle also responded.




Figure 7: keyboard_hri_node in a Rosject shell: start and pause keypresses published and confirmed, alongside the live Simulation panel.

Figure 8 is a live ros2 topic echo `/mission_state` capture. Oldest to newest, the capture shows AVOID_OBSTACLE, REPLAN, AVOID_OBSTACLE, IDLE, PLANNING, NAVIGATE and finally MISSION_COMPLETE. The code only goes from AVOID_OBSTACLE back to IDLE after ten failed avoidance attempts. So the mission most likely aborted, and I restarted it with `s`. It then planned again and reached MISSION_COMPLETE.




Figure 8: Live `/mission_state` echo showing the robot cycle through AVOID_OBSTACLE, REPLAN, IDLE, PLANNING, NAVIGATE and MISSION_COMPLETE.

Figure 9 is `monitor_node`'s summary from partway through the run, and Table 4 copies its values. The header still says "in progress", so this is not the final summary. The state times add up to 133.5 s against a total of 172.8 s. `monitor_node` only adds a state's time when the robot leaves that state, so the state running at the moment of capture is missing.




Figure 9: `monitor_node`'s mission summary, captured while the mission was in progress.

Table 4: Mission summary logged by monitor_node on TheConstruct.ai, Rosject XctZpczWOA, while the mission was in progress.

Metric	Value
Total mission time	172.8 s
Distance travelled	12.16 m
Obstacle encounters	11
Time in AVOID_OBSTACLE	10.3 s
Time in IDLE	61.7 s
Time in NAVIGATE	59.8 s
Time in PLANNING	0.0 s
Time in REPLAN	1.7 s

Table 5 lists the test cases that the captured evidence supports.

Table 5: Task 2 test cases and observed outcomes.

Test	Input	Observed outcome
Start from IDLE	s pressed	IDLE, PLANNING, NAVIGATE (Figures 7 and 8)
Obstacle on route	Pallets and pylons near the path	AVOID_OBSTACLE and REPLAN entered, 11 encounters logged (Figure 8, Table 4)
Repeated failed avoidance	Gap not cleared	Dropped to IDLE, most likely the ten-attempt abort, restarted with s (Figure 8)
Restart and full patrol	s pressed again	PLANNING, NAVIGATE, MISSION_COMPLETE (Figure 8)
Path toggle	v pressed	Planned-path line logged as ON (Figure 7)

The local Jazzy/Harmonic workspace gave one extra view that is hard to capture in the Rosject browser window. Figure 10 shows the LIDAR rays near a pallet obstacle with the forward camera feed beside them.




Figure 10: Local Gazebo Harmonic verification: LIDAR rays visualised around the robot near a pallet obstacle, with the forward camera feed shown alongside.

The Rosject evidence has five parts. These are the Rosject identity and live simulation (Figure 5), the running world (Figure 6), keyboard HRI with acknowledgement (Figure 7), a state sequence ending in MISSION_COMPLETE (Figure 8) and the logged metrics (Figure 9 and Table 4). The LIDAR view in Figure 10 comes from the local workspace.

4.4 Evaluation and Reflection

The main problem in development was over-cautious obstacle avoidance. The robot took wide detours around gaps it could have fitted through. The code comments show the cause. The margin was 0.09 m on top of the half-width, so the robot treated itself as having a half-width of about 0.20 m, or roughly 0.40 m wide. I cut it to 0.05 m. My first scoring rule also picked the heading closest to the goal even with little clearance, so the robot oscillated at gap edges. I changed it to prefer the heading with more open space among similar ones, and I retested after each change.

Table 4 shows eleven obstacle encounters over 12.16 m, with 10.3 s in AVOID_OBSTACLE against 59.8 s in NAVIGATE. The 61.7 s in IDLE most likely includes the wait before the first start and the abort in Figure 8.

REPLAN is also simple. It drives forward along a clear heading for 0.45 m and does not plan around the whole obstacle. A local costmap-based planner, as used within the Nav2 navigation

stack (Macenski et al., 2020), would route around larger obstacles better, but it is much more work. The pause key also has no resume, because s is only accepted from IDLE.

Explicit states, reactive avoidance with an odometry backup and constant logging are easy to explain to warehouse staff (Siegwart, Nourbakhsh and Scaramuzza, 2011). The figures in Table 4 are the kind an operator could check.

Combined Evaluation and Reflection

Task 1 gives every object in a static scene a role. I assigned the roles by hand, and the code retrieves them with dictionary lookups, as Table 1 and the nearest-target result in Section 3.5 show. In Task 2 the robot moves between six behaviour states, such as avoiding an obstacle, replanning and completing the mission. It reports which state it is in as it goes.

In both tasks a data step fed a decision step. In Task 1 the queries read object positions from PyBullet and semantic_db. In Task 2 the LIDAR and odometry readings drove the state machine. Both were easier to check because the system reported what it was doing. The most useful lesson came from Task 2. The avoidance was too cautious and took several rounds of threshold tuning to fix. A robot can avoid every obstacle and still behave badly, and the monitoring output helped me see that.

Concluding Remarks

Task 1 produced a PyBullet room with fourteen labelled objects and six camera views. The role lookups worked as intended, including the nearest-target query, which returned 0.092 m. Task 2 produced a TurtleBot3 patrol in ROS2 that reached MISSION_COMPLETE in Rosject XctZpczWOA. That run was not clean. The state log shows the robot dropping back to IDLE once, and I restarted it. Table 4 also comes from a single run, so I cannot say how repeatable the figures are. With more time I would record an uninterrupted run with the final monitor summary. I would also add a resume key, since pause currently has none. For Task 1 I would try predicting roles from the rendered images. The camera framing and the over-cautious avoidance each took several rounds of changes before the results were usable.

Acknowledgement

Perplexity AI (2026) was used to support literature searching, identify potentially relevant academic sources, and assist with grammar correction and word-limit reduction during editing.

All sources included in the reference list were checked by the author and cited as the original academic works rather than citing Perplexity AI as a source. The analysis, interpretation, discussion, and final written content submitted in this report are the author's own work and responsibility.

The architecture and pipeline diagrams in Sections 3.3 and 4.2 were generated using Python scripts with the matplotlib library, written specifically to render the structures described in this report. The scripts and their outputs were reviewed by the author before inclusion.

References

Borenstein, J. and Koren, Y. (1991) 'The vector field histogram - fast obstacle avoidance for mobile robots.' IEEE Transactions on Robotics and Automation 7(3): 278-288

Colledanchise, M. and Ögren, P. (2018) Behavior Trees in Robotics and AI: An Introduction. Boca Raton, FL: CRC Press

Collins, J., Chand, S., Vanderkop, A. and Howard, D. (2021) 'A Review of Physics Simulators for Robotic Applications.' IEEE Access 9: 51416-51431. [online] Available from <https://doi.org/10.1109/ACCESS.2021.3068769> [Accessed 20 September 2026]

Coumans, E. and Bai, Y. (2021) PyBullet, a Python Module for Physics Simulation for Games, Robotics and Machine Learning. [online] Available from <https://github.com/bulletphysics/bullet3> [Accessed 20 September 2026]

Kostavelis, I. and Gasteratos, A. (2015) 'Semantic mapping for mobile robotics tasks: A survey.' Robotics and Autonomous Systems 66: 86-103. [online] Available from <https://doi.org/10.1016/j.robot.2014.12.006> [Accessed 20 September 2026]

Macenski, S., Martin, F., White, R. and Ginés Clavero, J. (2020) 'The Marathon 2: A Navigation System.' In: Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS). Las Vegas, 24-30 October 2020: 2718-2725

McCormac, J., Handa, A., Leutenegger, S. and Davison, A.J. (2017) 'SceneNet RGB-D: Can 5M Synthetic Images Beat Generic ImageNet Pre-training on Indoor Segmentation?' In: Proceedings of the IEEE International Conference on Computer Vision (ICCV). Venice, 22-29 October 2017: 2678-2687

Open Robotics (2024) Gazebo Harmonic Documentation. [online] Available from <https://gazebosim.org/docs/harmonic> [Accessed 20 September 2026]

Open-Source Robotics Foundation (2024) ROS 2 Documentation: Humble. [online] Available from <https://docs.ros.org/en/humble/index.html> [Accessed 20 September 2026]

Perplexity AI (2026) Perplexity [AI tool]. [online] Available from <https://www.perplexity.ai> [Accessed 20 September 2026]

PyBullet (2023) PyBullet Quickstart Guide. [online] Available from <https://pybullet.org/wordpress/> [Accessed 20 September 2026]

Quigley, M., Gerkey, B. and Smart, W.D. (2015) Programming Robots with ROS: A Practical Introduction to the Robot Operating System. Sebastopol, CA: O'Reilly Media

ROBOTIS (2023) TurtleBot3 e-Manual. [online] Available from <https://emanual.robotis.com/docs/en/platform/turtlebot3/overview/> [Accessed 20 September 2026]

Siegwart, R., Nourbakhsh, I.R. and Scaramuzza, D. (2011) Introduction to Autonomous Mobile Robots. 2nd ed. Cambridge, MA: MIT Press

TheConstruct (2023) ROS2 Basics in 5 Days (Python). [online] Available from <https://www.theconstruct.ai/> [Accessed 20 September 2026]